

VGG CNN FOR URBAN SOUND TAGGING

Technical Report

Clément Gousseau

Orange Labs Lannion, France. clement.gousseau@gmail.com

ABSTRACT

A model of urban sound tagging is presented (Task 5 of DCASE 2019 [1][2]). The task is to detect activities from 10-seconds audio segments recorded in the streets of New York City (SONYC dataset). The model is based on the model presented in the book *Hands-On Transfer Learning with Python* [3] which does urban sound classification for the *UrbanSound* dataset. This model has been adapted and optimized to address the task 5 of DCASE2019. It achieved a AUPRC of 82.6 for the coarse-grained model where the baseline achieves an AUPRC of 76.2.

1. INTRODUCTION

The development dataset is composed of 2351 recordings in the training dataset and 443 recordings in the validation dataset. Each recording is a 10-seconds audio segment recorded in the streets of New York City. For the training set, each recording has been annotated by three volunteers on Zooniverse, a web platform for citizen science. For the validation set, each recording has been annotated by the organizers of the DCASE Challenge. The annotations follow a fined-grained or a coarse-grained taxonomy (see the figure below).

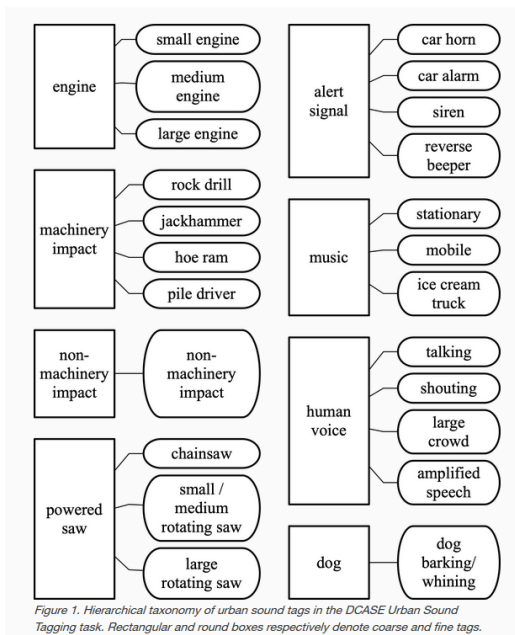


Figure 1: DCASE Task 5 Taxonomy

Labels are non-exclusive: several (or none) classes can be present in each recording. For each recording, the goal is to output a probability of presence for each class where 0 corresponds to an unlikely presence and 1 to a likely presence. The performance of the model is measured using micro-averaged AUPRC.

2. PROPOSED MODELS

2.1. Features

The feature engineering uses the method presented in the book *Hands-On Transfer Learning with Python*. First the recordings are re-sampled using a sampling rate of 22050Hz. Then three features are extracted from these signals:

- the mel-spectrograms using 64 mel-bands and a hop length of 512 thus resulting a 64 rows x 431 columns image
- the averaged value of the harmonic and percussive components (64 rows x 431 columns image)
- the derivative of the log-mel spectrograms (64 rows x 431 columns image)

This spectrograms computations has been done using the *librosa* library [4]. The figure below represents the three images extracted from an alert-signal recording.

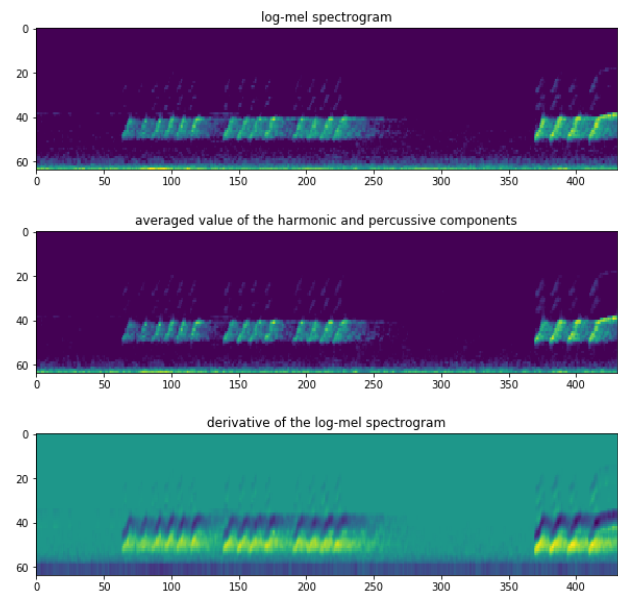


Figure 2: The three input images corresponding to a recording

2.2. Model structure

A VGG-style [5] convolutional neural network is used to detect the classes from the input images:

Input 64 x 431 x 3
64 x conv 3x3
64 x conv 3x3
MaxPooling 2x2
128 x conv 3x3
128 x conv 3x3
MaxPooling 2x2
256 x conv 3x3
256 x conv 3x3
256 x conv 3x3
MaxPooling 2x2
512 x conv 3x3
512 x conv 3x3
512 x conv 3x3
MaxPooling 2x2
512 x conv 3x3
512 x conv 3x3
512 x conv 3x3
MaxPooling 2x2
Flatten
1024-Fully Connected + L2-regularization
ReLu Activation
Dropout
1024-Fully Connected + L2-regularization
ReLu Activation
Dropout
512-Fully Connected + L2-regularization
ReLu Activation
Dropout
512-Fully Connected + L2-regularization
ReLu Activation
Dropout
n_classes-Fully Connected + L2-regularization
Sigmoid Activation

All convolutional layers are initialized with the weights of VGG16 pre-trained on Imagenet dataset but they remain unfrozen during the training.

$n_classes$ depends on the targets (fine-grained or coarse-grained labels).

2.3. Data augmentation

The training set is quite small (2351 samples). Data augmentation is a way to artificially increase the size of the training set and avoid overfitting. Mixup is a data augmentation method that has been experimented for this task. From two samples $\{input : x_1, target : y_1\}$ and $\{input : x_2, target : y_2\}$ a 'new' sample is created: $\{input : x_3 = \lambda x_1 + (1 - \lambda)x_2, target : y_3 = \lambda y_1 + (1 - \lambda)y_2\}$ where $\lambda \sim \beta(mixup_rate)$

2.4. Model optimization

Hyperparameter tuning is an important but tricky part of machine learning problems. For this task, Particle Swarm Optimization [6] has been used to optimize some hyperparameters of the network.

These hyperparameters are: *dropout_rate*, *L2-regularization constant*, *batch size*, *mixup_rate*.

Particle Swarm Optimization is a population-based optimization method, it was inspired from animal social groups like herds, schools and flocks. The swarm is composed by particles which have a position in the hyperparameter search space. Particles move in the search space and cooperate according to simple mathematical formulas in order to find an optimal solution. Each particle is driven by three components: an inertia component, an individual component (which drives each particle towards the best position the particle has explored so far), a collective component (which drives all the particles towards the best position the swarm has explored so far).




Figure 3: Schematic representation of updating the velocity of a particle [7]

Here, 5 particles were used for the optimization. 6 iterations were done, then 30 hyperparameters settings were evaluated. The hyperparameter setting that gave the best score was:

- batch size: 1
- constant for L2-regularization: 0.1
- mixup rate: 0.85
- dropout rate: 0.3

2.5. Model training

The model is trained using the following hyperparameters:

- loss: binary crossentropy as defined in the baseline
- optimizer: Adam
- learning rate: 1e-5
- batch size: 1
- constant for L2-regularization: 0.1
- mixup rate: 0.85
- dropout rate: 0.3
- number of epochs: 100

2.6. Model selection

The model is trained during 100 epochs and AUPRC is computed at each epoch. The best 4 models that show the best AUPRC are stored and the final prediction is the average of these 4 models.

3. RESULTS

3.1. Coarse-level model

	best pre-trained VGG	Baseline
Micro AUPRC	82.6	76.2
Micro F1-score	74.3	67.4
Macro AUPRC	61.1	54.2
Coarse tag AUPRC		
engine	86.8	85.5
machinery impact	60.5	36.0
non-machinery impact	56.5	36.1
powered saw	68.9	67.9
alert signal	92.1	81.3
music	18.0	29.9
human voice	94.8	94.5
dog	11.4	2.8

3.2. Fine-level model

3.2.1. Fine-level evaluation

	best pre-trained VGG	Baseline
Micro AUPRC	70.1	67.2
Micro F1-score	61.3	50.2
Macro AUPRC	47.2.8	42.7
Coarse tag AUPRC		
engine	65.3	71.2
machinery impact	25.1	19.8
non-machinery impact	54.3	36.4
powered saw	31.2	38.6
alert signal	83.3	63.6
music	11.8	21.5
human voice	88.7	88.0
dog	18.2	2.9

3.2.2. Coarse-level evaluation

	best pre-trained VGG	Baseline
Micro AUPRC	77.4	74.3
Micro F1-score	63.8	50.7
Macro AUPRC	56.7	53.0
Coarse tag AUPRC		
engine	85.2	85.9
machinery impact	27.0	28.5
non-machinery impact	54.3	36.4
powered saw	64.2	72.0
alert signal	90.4	75.3
music	18.7	28.3
human voice	95.7	94.3
dog	18.2	2.9

4. CONCLUSION

Tagging of urban sounds was investigated. A model based on pre-trained VGG-style network was developed and submitted to the challenge DCASE 2019 (task 5). The best model achieves an AUPRC of 82.6 for the coarse-level prediction and 70.1 for the fine-level prediction.

5. REFERENCES

- [1] <http://dcase.community/challenge2019/>.
- [2] J. P. Bello, C. Silva, O. Nov, R. L. Dubois, A. Arora, J. Salamon, C. Mydlarz, and H. Doraiswamy, "Sonyc: A system for monitoring, analyzing, and mitigating urban noise pollution," *Communications of the ACM*, vol. 62, no. 2, pp. 68–77, Feb 2019.
- [3] D. Sarkar, R. Bali, and T. Ghosh, *Hands-On Transfer Learning with Python: Implement advanced deep learning and neural network models using TensorFlow and Keras*. Packt Publishing, 2018. [Online]. Available: <https://books.google.fr/books?id=aPFsDwAAQBAJ>
- [4] B. McFee, M. McVicar, S. Balke, V. Lostanlen, C. Thom, C. Raffel, D. Lee, K. Lee, O. Nieto, F. Zalkow, D. Ellis, E. Battenberg, R. Yamamoto, J. Moore, Z. Wei, R. Bittner, K. Choi, nullmightybofo, P. Friesch, F.-R. Stter, Thassilo, M. Vollrath, S. K. Golu, nehz, S. Waloschek, Seth, R. Naktinis, D. Repetto, C. F. Hawthorne, and C. Carr, "librosa/librosa: 0.6.3," Feb. 2019. [Online]. Available: <https://doi.org/10.5281/zenodo.2564164>
- [5] K. Simonyan and A. Zisserman, "Very deep convolutional networks for large-scale image recognition," in *International Conference on Learning Representations*, 2015.
- [6] Y. Shi and E. RC, "A modified particle swarm optimizer," vol. 6, 06 1998, pp. 69 – 73.
- [7] A. Ahmadi, F. Karray, and M. S. Kamel, "Flocking based approach for data clustering," *Natural Computing*, vol. 9, pp. 767–791, 09 2010.